

LLMs Can Predict Failure Risk, But Struggle to Predict Which Collaboration Protocol Pays Off: Cost-Aware Protocol Routing Across Reasoning Tasks

Chih-Hsuan Yang¹, Jingyan Jiang¹, Cheng-Hau Yang¹, Vikram Vasudevan²,
Huihuo Zheng¹, Venkatram Vishwanath¹, Rajeep Thakur¹

¹Argonne National Laboratory, Lemont, IL, USA

²Oregon State University, Corvallis, OR, USA

bellayang@anl.gov

Abstract

Multi-agent large language model (LLM) systems can improve reasoning by spending more computation, but deployment requires deciding when extra collaboration is worth its cost. We isolate this decision by running every problem under four protocols while holding the solver fixed within each setting: direct solving (Baseline), iterative self-correction (Single), planner-executor-reviewer collaboration (PER), and multi-agent deliberation (Broadcast). The primary benchmark comprises 4,181 competition-level math problems; paired robustness checks cover four benchmarks spanning competition math, biology, and broader science with two solver families. Across fixed policies, trained routers, and frozen LLM routers, conservative policies under-escalate, whereas higher-solve frozen routers often over-escalate. A post-answer, pre-collaboration gpt-oss-120b probe ranks Baseline failures with 0.8847 AUROC (4,151 parseable cases; 95% CI [0.8732, 0.8955]). The same score remains informative for predicting whether any collaboration helps (0.7683 AUPRC), but is much weaker for identifying PER- or Broadcast-specific value (0.1674 and 0.1041 AUPRC). Separately, the pre-answer self-confidence gate reaches 78.0% solve at 45K tokens, compared with 73.8% at 71.3K for a frozen gpt-oss-120b router and 92.4% for a retrospective fixed-order oracle. Across 10 paired model-condition settings, the oracle adds 23.2–58.3 points of retrospective coverage over Baseline, but protocol profiles vary by task. In the six settings with held-out router evaluations, oracle gaps remain 18.5–28.9 points. Confidence can therefore support initial escalation, while protocol-specific cost-aware routing remains unresolved.

1 Introduction

Large language model (LLM) reasoning systems can spend drastically different amounts of computation on the same problem. A direct solver may cost tens of thousands of tokens, while self-

correction or multi-agent deliberation can cost an order of magnitude more. Stronger protocols are often more accurate, but deploying the strongest protocol everywhere silently multiplies cost. Deployment therefore requires deciding both *whether* to escalate and *which* collaboration protocol is worth its marginal cost.

We build a matched benchmark of 4,181 competition-level math problems, each evaluated under four protocols with sharply different token costs and solve rates. Observing every protocol on every problem permits offline evaluation against realized solve and cost outcomes and a retrospective fixed-order oracle. This design reveals a directional failure hidden by routing accuracy: cost-conservative policies under-escalate, whereas higher-solve frozen LLM routers often over-escalate.

These observations motivate a sharper distinction. *Failure-risk prediction* asks whether Baseline will be wrong; *collaboration-value prediction* asks whether, and which, stronger protocol justifies its added cost. A full-benchmark post-answer probe directly supports the first claim, but its precision degrades for increasingly protocol-specific value targets. Matched outcome checks extend across four benchmarks spanning competition math, biology, and broader science, using both gpt-oss-120b and Gemma-4-31B-it (Gemma Team, 2026; Laurent et al., 2024; Arora et al., 2023; Wang et al., 2024). The more complete confidence and held-out-router evaluations remain a targeted six-setting subset, so this breadth is evidence of task dependence rather than universality.

Several related literatures vary different inference components. Model routing selects among models or cascades with different capabilities and prices (Chen et al., 2024; Ong et al., 2025; Somerstep et al., 2025; Song et al., 2025). Adaptive computation and sample routing vary reasoning depth or the number of sampled paths (Wang et al., 2023;

Balachandran et al., 2025; Wu et al., 2025). Tool and multi-agent routing vary tools, roles, or collaboration structures (Yue et al., 2025; Gan et al., 2025; Zhu et al., 2025). Calibration work instead asks whether confidence tracks correctness (Joshi et al., 2025; Li et al., 2025a,b). Our setting holds the solver family fixed within each comparison and changes only the collaboration protocol. This isolates protocol value from a change in base model capability; Appendix Table S1 summarizes these distinctions.

Contributions. First, we provide matched four-protocol outcomes and solve-cost evaluation that keep the solver fixed within each setting, with paired coverage checks across four benchmarks. Second, we identify asymmetric under- versus over-escalation across heuristic, learned, and frozen-LLM routers. Third, we empirically separate failure risk from protocol-specific collaboration value: same-model confidence supports an initial stay-or-escalate decision, while substantial fixed-order-oracle gaps and weak protocol-specific precision leave full cost-aware routing open. The source package includes machine-readable aggregate tables, and the companion dataset archive contains anonymized protocol traces and outcome labels for the matched comparisons.

2 Task, Benchmark, and Metrics

Protocol routing. A *router* selects one action before observing protocol outcomes. Fixed policies are degenerate routers; learned and frozen-LLM routers map problem text and/or allowed metadata to BASELINE, SINGLE, PER, BROADCAST, or None. The None action abstains and can be both an oracle label and a router prediction.

Our main benchmark uses the clean exact-answer Omni-MATH 2 subset, a manually revised release derived from Omni-MATH, with 4,181 competition-level math problems (Gao et al., 2024; Ballon et al., 2026). Router-visible metadata includes source, domain path, numeric difficulty, and a ten-level difficulty tier; real examples include source values cayley, fermat, and pascal; one domain path is Mathematics $\rightarrow$ Algebra $\rightarrow$ Prealgebra $\rightarrow$ Simple Equations. For the LAB-Bench robustness study, visible fields are dataset name, domain, slice/subset, original identifier, prompt condition, and subtask, e.g., Biology $\rightarrow$ LAB-Bench $\rightarrow$ CloningScenarios. Routers and confidence probes never receive gold answers,

correctness labels, oracle labels, or protocol outcomes.

Protocols. In the main benchmark, every problem is run once under four protocols with the same gpt-oss-120b solver stack (OpenAI, 2025). BASELINE is one direct attempt without self-correction. SINGLE adds iterative self-correction. PER uses planner, executor, and reviewer roles. BROADCAST uses multi-agent deliberation with shared candidates and peer approval. On the main held-out split, these protocols average 18.2K, 47.6K, 401.9K, and 622.1K tokens and solve 56.3%, 78.5%, 84.9%, and 88.9%, respectively. Runs use temperature 0.0; protocol definitions and compute accounting are detailed in Appendix C.3.

Fixed-order oracle. For each problem, the retrospective oracle is the first successful action in the aggregate cost order $\text{BASELINE} < \text{SINGLE} < \text{PER} < \text{BROADCAST}$; it is None if all four fail. This is a matched, single-realization diagnostic upper bound, not a deployable policy, a per-instance minimum-token oracle, or expected success under repeated sampling. On the full benchmark, its labels are 56.8% Baseline, 23.0% Single, 8.8% PER, 4.2% Broadcast, and 7.3% None.

Routers and evaluation. The main learned-router comparison uses a stratified 80/10/10 split (seed 42; 3,342/416/423 problems). *Tier-majority* predicts each difficulty tier’s train-split majority fixed-order-oracle label, falling back to the train-split global majority. It is a transparent metadata-only sanity check, not a deployed heuristic. Learned routers are five-class logistic regressions over allowed metadata, with optional TF-IDF word uni- and bigrams from the problem text; model and hyperparameter selection use the development split only. Frozen routers receive the same text and metadata in a label-only prompt. The cost-aware ablation also provides numeric average protocol costs and short routing examples.

We report solve rate, average tokens, and *excess tokens*, the mean positive per-problem token overpayment relative to the realized oracle. Under- and over-escalation mean choosing cheaper and costlier actions than the oracle. None has zero protocol cost and ranks below Baseline; router or probe overhead still counts. Hence, choosing a protocol when the oracle is None is over-escalation, and all its tokens are excess.




Figure 1: (a) Fixed-order-oracle labels by difficulty tier; fractions sum to 1 within each tier. (b) Solve-cost frontier relative to Tier-majority; upper-left is better. The self-confidence gate improves the low-cost frontier, while larger frozen routers buy additional solves through substantially more token use. Gemma-4-E4B here is a frozen router over the main gpt-oss traces, not the Gemma-4-31B-it solver used in Section 5.

Class	Policy	Solve (%) [95% CI]	Avg. tokens (K) [95% CI]	Excess (K) [95% CI]
Fixed	Baseline	56.3 [51.5, 61.0]	18.2 [17.1, 19.2]	1.7 [1.1, 2.4]
Heuristic	Tier-majority	65.0 [60.5, 69.5]	28.9 [26.0, 31.9]	5.7 [3.9, 7.7]
Frozen LLM	gpt-oss-120b	73.8 [69.5, 77.5]	71.3 [56.5, 86.7]	37.1 [24.2, 50.6]
Frozen LLM	gpt-oss + cost prompt	78.3 [74.2, 82.0]	88.6 [71.0, 107.6]	51.6 [35.2, 68.9]
Confidence	Self-confidence gate	78.0 [74.0, 81.8]	45.0 [41.6, 48.7]	14.8 [12.2, 17.4]
Retrospective	Fixed-order oracle	92.4 [89.8, 95.0]	101.1 [80.1, 124.5]	0.0 [0.0, 0.0]

Table 1: Representative routing policies on the primary held-out test split ($n = 423$). Intervals use 2,000 problem-level bootstrap resamples. Excess is positive token overpayment above the realized fixed-order oracle. The protocol executions and router/probe calls use deterministic decoding, so the intervals do not measure fresh-run variability.

3 Routing Errors and Cost Tradeoffs

We compare fixed and heuristic policies, lightweight trained routers, frozen LLM routers, a cost-aware prompt, and the confidence policy introduced in Section 4. Figure 1 shows the oracle composition by tier and the resulting solve-cost frontier. Table 1 reports representative policies with uncertainty; the full grouped comparison appears in Appendix Table S9.

Cheap references are competitive. Baseline solves 56.3% at 18.2K average tokens. Tier-majority raises solve rate to 65.0% at 28.9K tokens. The selected metadata-only and text+metadata logistic routers solve 60.8% and 61.2%, respectively, below this simple reference. Balanced and embedding variants recover more solves only by routing more often to expensive actions; full results and feature ablations appear in Appendix B.

Frozen routers trade cost for fewer misses. The frozen gpt-oss-120b router reaches 73.8% solve at 71.3K tokens. Showing numeric protocol costs and routing examples raises it to 78.3%, but also to 88.6K tokens and 51.6K excess. Larger cross-family frozen routers reach 80–83% solve while spending 156K–280K tokens on average. No evaluated router approaches the fixed-order-oracle operating point.

Errors are directional. Tier-majority under-escalates on 27.4% and over-escalates on 12.5% of test problems. The gpt-oss-120b router cuts under-escalation to 18.0% but raises over-escalation to 33.3%. Higher-solve Llama and Gemma frozen routers reduce under-escalation to 6–11% while over-escalating on 63–71%. Thus, router gains are not interchangeable: conservative policies miss recoverable problems; aggressive routers buy solves through costly escalation.

4 Failure Risk Is Not Protocol Value

The directional errors raise two different questions. Can the solver recognize that its Baseline answer is likely wrong? If so, can it identify which stronger protocol is worth the added cost?

Post-answer failure-risk probe. After gpt-oss-120b produces its Baseline answer, we ask for $P(\text{Baseline correct})$. The probe sees only the problem, allowed metadata, and its own Baseline *final answer*; it sees no reasoning trace, gold answer, correctness label, oracle label, or collaboration outcome. Of 4,181 matched problems, 4,151 produce parseable scores (99.28%). The 30 unparseable outputs are excluded rather than imputed. Using $1 - P(\text{correct})$ as failure risk yields 0.8847 AUROC (95% CI [0.8732, 0.8955]), 0.8950 AUPRC, and 0.0852 expected calibration error.

This differs from the *pre-answer confidence probe*, which asks for single-pass solve probability before the model sees a Baseline answer. That probe has 0.859 AUROC on 329 usable estimates from the primary 423-example test split. The two results correspond to distinct operating points and are not treated as replicates.

Protocol-specific precision is much weaker. Table 2 applies the same no-leakage failure score to increasingly specific collaboration-value targets. It remains useful for the coarse question of whether any protocol improves on Baseline. For PER and Broadcast, however, AUPRC falls to 0.1674 and 0.1041. These exceed the rare target prevalences, so confidence is not devoid of value signal, but it is not a reliable full protocol selector.

A binary gate is useful but narrower. The self-confidence gate uses the earlier pre-answer confidence score: keep Baseline at confidence ≥ 70 , otherwise choose Single; missing confidence scores

Target	Prev.	AUROC	AUPRC
Baseline fails	43.4	0.8847	0.8950
Any collaboration helps	36.0	0.8544	0.7683
PER first success	8.8	0.7259	0.1674
Broadcast-only success	4.2	0.7639	0.1041

Table 2: The same post-answer failure score on increasingly specific targets ($n = 4,151$ parseable cases). Prevalence is a percentage. “PER first” means Baseline and Single fail but PER succeeds; “Broadcast-only success” means Baseline, Single, and PER fail while Broadcast succeeds in the matched realized runs.

also choose Single. The threshold is selected on dev by a knee-point rule (Satopää et al., 2011; Deb and Gupta, 2011). On the full test split it reaches 78.0% solve at 45.0K tokens, versus 73.8% at 71.3K for the frozen gpt-oss router (Table 1). The overlapping solve intervals support a cost-efficiency comparison, not a claim of higher solve rate. We reserve *self-confidence gate* for this binary policy; *two-threshold cascade* denotes the separate ablation in Appendix B, which uses Baseline, Single, and Tier-majority with high/low thresholds of 70/10.

The gate does not choose PER, Broadcast, or None. Its 14.4-point gap to the fixed-order oracle and the target-specific results identify the remaining problem: estimating the marginal value of each expensive protocol after the initial escalation decision.

5 Matched Breadth and Router Robustness

To test whether the matched protocol pattern is confined to OmniMath, we extend the analysis to four benchmarks across three broad task families: OmniMath 2 (competition math), JEEBench (engineering entrance-exam STEM), SciBench (college-level science), and LAB-Bench (biology multiple-choice). LAB-Bench contributes two no-tool prompt conditions: “strict” includes relevant evidence in the prompt, while “text-no-tool” is a broader text-only slice. Both gpt-oss-120b and Gemma-4-31B-it cover all five conditions, yielding 10 paired settings.

Stronger collaboration has recoverable value in all 10 paired settings (Figure 2): fixed-order-oracle coverage exceeds Baseline by 23.2–58.3 points. The protocol profile varies, however. Broadcast is the strongest fixed protocol in nine settings, while PER exceeds Broadcast by 2.7 points for Gemma on SciBench. This breadth supports task-dependent protocol value, not universal generalization.

The deeper confidence and held-out-router analyses use the narrower six-setting subset comprising both solvers on OmniMath and the two LAB-Bench conditions. We train the same pre-specified text+metadata logistic router using a stratified 70/15/15 split (seed 20260712), dev selection only, and refitting on train+dev. It improves over Baseline by 7.2–37.5 points in all six settings, yet remains 18.5–28.9 points below the fixed-order oracle. Relative to Tier-majority, gains are mixed.

	(a) gpt-oss-120b					(b) Gemma-4-31B-it				
	Base	Single	PER	Broad.	Oracle	Base	Single	PER	Broad.	Oracle
OmniMath 2 (n=4,181)	56.8	78.8	85.2	89.2	92.7	69.4	85.9	90.7	93.0	95.7
JEEBench (n=515)	41.5	55.3	91.5	95.0	96.5	70.5	81.4	95.9	98.2	99.2
SciBench (n=565)	62.1	72.2	87.4	89.7	91.2	70.3	79.3	91.3	88.7	93.5
LAB strict (n=741)	19.0	30.2	47.6	74.2	77.3	44.4	58.2	69.4	89.3	90.5
LAB text-no-tool (n=1,542)	30.5	55.6	59.3	81.3	86.4	42.0	74.8	85.0	91.0	96.4

Figure 2: Matched solve/coverage percentages across four benchmarks (five prompt conditions), four protocols, and two solver families. Cell hue identifies the protocol, and deeper tint indicates higher solve/coverage. The outlined cell is the strongest deployable fixed protocol in each row. Oracle is retrospective fixed-order coverage computed from the same four realized outcomes. Unlike Table 1, which reports the primary held-out test split, this figure uses all matched problems available in each setting. Each value is one deterministic realized execution per problem–protocol pair.

Thus, observable text and metadata contain routing signal but do not yield uniformly reliable protocol-value routing.

The expensive protocols are also not interchangeable. Conditional on Baseline and Single both failing, Broadcast exceeds PER by 9.6–10.7 points on OmniMath and 30.8–44.6 points on LAB-Bench. Nevertheless, PER-only successes remain in every setting (2.4–11.1% of these conditional subsets), so Broadcast does not pointwise dominate PER.

6 Discussion and Conclusion

The matched design separates two decisions that aggregate routing accuracy can hide. Failure-risk estimation is useful for deciding whether to remain with a cheap direct solver. Collaboration-value estimation must additionally predict which expensive workflow provides enough marginal benefit. Conservative and aggressive routers fail in opposite directions, the binary gate covers only the first decision, and the broader held-out routers retain large oracle gaps.

Our conclusion is therefore diagnostic rather than a claim to have solved full routing. Same-model confidence can support an initial stay-or-escalate gate, but selecting among Single, PER, Broadcast, and None remains a cost-aware control problem. The benchmark, directional error analysis, and explicit failure-risk/protocol-value distinction provide measurable targets for that next step.

7 Limitations

Our main router comparison and pre-answer confidence gate use one solver family and one math benchmark. The broader paired matched-outcome study covers four benchmarks across three broad task families and two solvers. Post-answer confidence and held-out router analyses cover the narrower six settings formed by both solvers on OmniMath and the two LAB-Bench conditions. These are targeted robustness checks, not a claim of broad universality. Each problem–protocol pair has one deterministic realized outcome. The 2,000-resample intervals quantify variation over benchmark problems, not fresh-run stochastic variability. The fixed-order oracle is retrospective and uses an aggregate cost order; it is neither deployable nor an estimate of per-instance expected utility.

We measure cost primarily in logged tokens. Latency, monetary price, energy, parallelism, and answer quality beyond exact correctness can change deployment preferences. The post-answer probe excludes 30 unparseable outputs from its probability metrics, and confidence quality varies substantially by model and domain. Finally, the observed PER–Broadcast differences do not identify a causal mechanism; understanding why protocol value changes by task remains open.

Acknowledgments

This research used resources of the Argonne Leadership Computing Facility, a U.S. Department of Energy (DOE) Office of Science user facility at Argonne National Laboratory (ANL) operated under Contract No. DE-AC02-06CH11357.

References

- Daman Arora, Himanshu Singh, and Mausam. 2023. [Have LLMs advanced enough? a challenging problem solving benchmark for large language models](#). In *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing*, pages 7527–7543, Singapore. Association for Computational Linguistics.
- Vidhisha Balachandran, Jingya Chen, Lingjiao Chen, Shivam Garg, Neel Joshi, Yash Lara, John Langford, Besmira Nushi, Vibhav Vineet, Yue Wu, and Safoora Yousefi. 2025. [Inference-time scaling for complex tasks: Where we stand and what lies ahead](#). *arXiv preprint arXiv:2504.00294*.
- Marthe Ballon, Andres Algaba, Brecht Verbeken, and Vincent Ginis. 2026. [Benchmarks saturate when the model gets smarter than the judge](#). *arXiv preprint arXiv:2601.19532*.
- Lingjiao Chen, Matei Zaharia, and James Zou. 2024. [FrugalGPT: How to use large language models while reducing cost and improving performance](#). *Transactions on Machine Learning Research*.
- Kalyanmoy Deb and Shivam Gupta. 2011. [Understanding knee points in bicriteria problems and their implications as preferred solution principles](#). *Engineering Optimization*, 43(11):1175–1204.
- Bingzheng Gan, Yufan Zhao, Tianyi Zhang, Jing Huang, Li Yusu, Shu Xian Teo, Changwang Zhang, and Wei Shi. 2025. [Master: A multi-agent system with LLM specialized MCTS](#). In *Proceedings of the 2025 Conference of the Nations of the Americas Chapter of the Association for Computational Linguistics: Human Language Technologies (Volume 1: Long Papers)*, pages 9409–9426, Albuquerque, New Mexico. Association for Computational Linguistics.
- Bofei Gao, Feifan Song, Zhe Yang, Zefan Cai, Yibo Miao, Qingxiu Dong, Lei Li, Chenghao Ma, Liang Chen, Runxin Xu, Zhengyang Tang, Benyou Wang, Daoguang Zhan, Shanghaoran Quan, Ge Zhang, Lei Sha, Yichang Zhang, Xuancheng Ren, Tianyu Liu, and Baobao Chang. 2024. [Omni-MATH: A universal olympiad level mathematic benchmark for large language models](#). *arXiv preprint arXiv:2410.07985*.
- Gemma Team. 2026. [Gemma 4 technical report](#). *arXiv preprint arXiv:2607.02770*.
- Abhinav Joshi, Areeb Ahmad, and Ashutosh Modi. 2025. [Calibration across layers: Understanding calibration evolution in LLMs](#). In *Proceedings of the 2025 Conference on Empirical Methods in Natural Language Processing*, pages 14686–14714, Suzhou, China. Association for Computational Linguistics.
- Jon M. Laurent, Joseph D. Janizek, Michael Ruzo, Michaela M. Hinks, Michael J. Hammerling, Siddharth Narayanan, Manvitha Ponnampati, Andrew D. White, and Samuel G. Rodrigues. 2024. [LAB-Bench: Measuring capabilities of language models for biology research](#). *arXiv preprint arXiv:2407.10362*.
- Chengzu Li, Han Zhou, Goran Glavaš, Anna Korhonen, and Ivan Vulić. 2025a. [Large language models are miscalibrated in-context learners](#). In *Findings of the Association for Computational Linguistics: ACL 2025*, pages 11575–11596, Vienna, Austria. Association for Computational Linguistics.
- Yingjie Li, Yun Luo, Xiaotian Xie, and Yue Zhang. 2025b. [Task calibration: Calibrating large language models on inference tasks](#). In *Findings of the Association for Computational Linguistics: ACL 2025*, pages 6937–6951, Vienna, Austria. Association for Computational Linguistics.
- Isaac Ong, Amjad Almahairi, Vincent Wu, Wei-Lin Chiang, Tianhao Wu, Joseph E. Gonzalez, M. Waleed Kadous, and Ion Stoica. 2025. [RouteLLM: Learning to route LLMs with preference data](#). In *International Conference on Learning Representations*.
- OpenAI. 2025. [gpt-oss-120b & gpt-oss-20b Model Card](#). Accessed August 24, 2026.
- Ville Satopää, Jeannie Albrecht, David Irwin, and Barath Raghavan. 2011. [Finding a “kneedle” in a haystack: Detecting knee points in system behavior](#). In *2011 31st International Conference on Distributed Computing Systems Workshops*, pages 166–171. IEEE.
- Seamus Somerstep, Felipe Maia Polo, Allysson Flavio Melo de Oliveira, Prattyush Mangal, Mírian Silva, Onkar Bhardwaj, Mikhail Yurochkin, and Subha Maity. 2025. [CARROT: A cost aware rate optimal router](#). *arXiv preprint arXiv:2502.03261*.
- Wei Song, Zhenya Huang, Cheng Cheng, Weibo Gao, Bihan Xu, GuanHao Zhao, Fei Wang, and Runze Wu. 2025. [IRT-router: Effective and interpretable multi-LLM routing via item response theory](#). In *Proceedings of the 63rd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 15629–15644, Vienna, Austria. Association for Computational Linguistics.
- Xiaoxuan Wang, Ziniu Hu, Pan Lu, Yanqiao Zhu, Jieyu Zhang, Satyen Subramaniam, Arjun R. Loomba, Shichang Zhang, Yizhou Sun, and Wei Wang. 2024. [SciBench: Evaluating college-level scientific problem-solving abilities of large language models](#). In *Proceedings of the Forty-First International Conference on Machine Learning*, volume 235 of

Proceedings of Machine Learning Research, pages 50622–50649. PMLR.

Xuezhi Wang, Jason Wei, Dale Schuurmans, Quoc V. Le, Ed H. Chi, Sharan Narang, Aakanksha Chowdhery, and Denny Zhou. 2023. [Self-consistency improves chain of thought reasoning in language models](#). In *International Conference on Learning Representations*.

Menghua Wu, Cai Zhou, Stephen Bates, and Tommi Jaakkola. 2025. [Thought calibration: Efficient and confident test-time scaling](#). In *Proceedings of the 2025 Conference on Empirical Methods in Natural Language Processing*, pages 14291–14305, Suzhou, China. Association for Computational Linguistics.

Yanwei Yue, Guibin Zhang, Boyang Liu, Guancheng Wan, Kun Wang, Dawei Cheng, and Yiyang Qi. 2025. [MasRouter: Learning to route LLMs for multi-agent systems](#). In *Proceedings of the 63rd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 15549–15572, Vienna, Austria. Association for Computational Linguistics.

Kunlun Zhu, Hongyi Du, Zhaochen Hong, Xiaocheng Yang, Shuyi Guo, Zhe Wang, Zhenhailong Wang, Cheng Qian, Xiangru Tang, Heng Ji, and Jiaxuan You. 2025. [MultiAgentBench : Evaluating the collaboration and competition of LLM agents](#). In *Proceedings of the 63rd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 8580–8622, Vienna, Austria. Association for Computational Linguistics.

A Cross-Setting Robustness and Diagnostics

This section establishes the scope of the main claims through targeted robustness and diagnostic analyses. Unless stated otherwise, confidence intervals use 2,000 percentile bootstrap resamples over problem identifiers.

A.1 Positioning Within Routing and Adaptive Inference

Table S1 makes explicit which inference component each nearby line of work changes. Model routing and cascades select among base models (Chen et al., 2024; Ong et al., 2025; Somerstep et al., 2025; Song et al., 2025); adaptive computation and self-consistency allocate steps or samples (Wang et al., 2023; Balachandran et al., 2025; Wu et al., 2025); tool and agent routing changes tools, roles, or collaboration structures (Yue et al., 2025; Gan et al., 2025). Our matched analysis changes the collaboration protocol while holding the solver fixed within a setting and observing every action’s realized outcome.

A.2 Matched Protocol Coverage Across Four Benchmarks

We evaluate matched four-protocol outcomes on four benchmarks across three broad task families: OmniMath 2 (competition math), JEEBench (engineering entrance-exam STEM), SciBench (college-level science), and LAB-Bench (biology multiple-choice) (Ballon et al., 2026; Arora et al., 2023; Wang et al., 2024; Laurent et al., 2024). LAB strict is an LLM-only, no-tool/no-retrieval condition with relevant evidence included in the prompt; LAB text-no-tool is a broader text-only no-tool condition. Both openai/gpt-oss-120b (OpenAI, 2025) and google/gemma-4-31B-it (Gemma Team, 2026) cover all five evaluation conditions. Tables S2 and S3 therefore report 10 matched model-condition settings.

A.3 Post-Answer Confidence Probes

The model-call file and scoring labels are physically separated. The model reads probe_inputs.jsonl, containing problem text, allowed metadata, and its own Baseline final answer. Correctness and protocol outcomes are stored only in labels_for_scoring.csv and joined after inference. A schema validator rejects forbidden keys, and the saved leakage checks pass

in all six settings. No probe receives a reasoning trace.

For the primary gpt-oss OmniMath setting, 4,151 of 4,181 outputs are parseable. Its AUROC interval uses seed 20260712; the remaining tables use seed 20260714. Both use 2,000 problem-identifier resamples. The result differs from the 0.859 AUROC of the pre-answer confidence probe because that run uses 329 cleaned estimates from the 423-example test split and does not see a Baseline answer. The post-answer run uses the full matched setting, includes the Baseline final answer, allows repair attempts, and reaches 99.28% coverage.

A.4 Held-Out Text+Metadata Routers

For each setting, we stratify 70/15/15 by fixed-order-oracle label using seed 20260712. The pre-specified router family is five-class logistic regression over TF-IDF word unigrams/bigrams ($\text{min_df}=2$, at most 20,000 features) plus a sparse encoding of allowed metadata. The metadata features include setting/model identifiers and available dataset name, source, domain, difficulty/tier, subject, question type, subset, subtask, and prompt condition. We search $C \in \{0.25, 1, 4\}$ with or without balanced class weights. Selection uses dev macro-F1, then dev solve rate, then oracle gap; the selected configuration is refit on train+dev. No test outcome participates in feature or hyperparameter selection.

Protocol-token fields are complete for the gpt-oss outcome files but not for Gemma, so Gemma router costs are not estimated. For gpt-oss, average/excess tokens are 151.7K/105.5K (OmniMath), 556.4K/261.3K (LAB strict), and 266.0K/112.5K (LAB text-no-tool), using the same cost definitions as the main paper. Missing Gemma costs are reported as unavailable rather than imputed.

A.5 PER-Broadcast Interaction

Table S8 focuses on the decision point where both cheaper protocols fail. Broadcast is stronger on average, but PER-only successes occur in every setting. Thus, a global ordering does not imply pointwise dominance. The larger LAB differences are an observed domain dependence; we do not have direct measurements establishing a causal explanation.

For completeness, in the six confidence/router settings ordered as Gemma OmniMath, gpt-oss OmniMath, Gemma LAB strict, gpt-oss LAB strict, Gemma LAB text-no-tool, and gpt-oss LAB text-

Family	Routed unit	Fixed solver?	Matched outcomes?	Main target
Model routing/cascades	Base model	No	Usually no	Quality–price tradeoff
Adaptive computation	Steps/depth	Usually	Usually no	Stop or continue
Self-consistency routing	Sampled reasoning paths	Yes	Usually no	Number of samples
Tool/agent routing	Tools, roles, topology	Varies	Usually no	Workflow selection
This work	Collaboration protocol	Yes	Yes	Realized solve–token value

Table S1: Scope comparison. “Matched” means that every candidate action is observed for every evaluated problem, enabling paired offline analysis.

Setting	<i>n</i>	Baseline	Single	PER	Broadcast	Oracle
<i>gpt-oss-120b</i>						
OmniMath 2	4,181	56.83 [55.4,58.4]	78.78 [77.6,80.0]	85.17 [84.1,86.2]	89.21 [88.3,90.2]	92.68 [91.9,93.5]
JEEBench	515	41.55 [37.5,45.8]	55.34 [51.1,59.6]	91.46 [88.9,94.0]	94.95 [93.0,96.7]	96.50 [94.8,98.1]
SciBench	565	62.12 [58.1,66.0]	72.21 [68.5,75.8]	87.43 [84.6,90.1]	89.73 [87.1,92.2]	91.15 [88.7,93.5]
LAB strict	741	19.03 [16.1,21.9]	30.23 [27.0,33.7]	47.64 [44.0,51.3]	74.22 [71.1,77.5]	77.33 [74.2,80.3]
LAB text-no-tool	1,542	30.54 [28.3,32.9]	55.64 [53.2,58.1]	59.27 [56.7,61.7]	81.26 [79.4,83.2]	86.38 [84.6,88.1]
<i>Gemma-4-31B-it</i>						
OmniMath 2	4,181	69.39 [68.0,70.8]	85.94 [84.9,87.0]	90.65 [89.8,91.5]	92.99 [92.2,93.8]	95.67 [95.0,96.3]
JEEBench	515	70.49 [66.6,74.4]	81.36 [78.1,84.7]	95.92 [94.2,97.5]	98.25 [96.9,99.2]	99.22 [98.5,99.8]
SciBench	565	70.27 [66.6,74.0]	79.29 [75.8,82.5]	91.33 [89.0,93.5]	88.67 [85.8,91.2]	93.45 [91.3,95.4]
LAB strict	741	44.40 [40.9,47.9]	58.16 [54.5,61.7]	69.37 [66.0,72.6]	89.34 [87.0,91.5]	90.55 [88.4,92.6]
LAB text-no-tool	1,542	41.96 [39.5,44.4]	74.77 [72.6,76.9]	85.02 [83.3,86.8]	91.05 [89.6,92.5]	96.37 [95.4,97.3]

Table S2: Solve/coverage percentages with 95% problem-level bootstrap intervals across all 10 matched settings. Point estimates show two decimal places and interval endpoints one; machine-readable results retain full precision. Each point estimate is based on one deterministic realized run per problem–protocol pair.

no-tool, Broadcast-minus-PER on all problems is 2.3, 4.0, 20.0, 26.6, 6.0, and 22.0 points. Conditional on Baseline failure, the gaps are 6.3, 8.5, 35.0, 30.8, 10.5, and 30.3 points. On parseable post-answer confidence scores below 70, they are 5.0, 8.8, 25.3, 29.7, 7.8, and 21.9 points. The primary conditional comparison and its paired intervals are also included in the ancillary result table `per_broadcast_interaction.csv`.

A.6 Post-Answer Probe Prompt and Parsing

The default prompt is reproduced below. Runs use temperature 0.0, a 1,024 token cap for the gpt-oss OmniMath post-answer probe, at most four API retries, and up to two JSON repair attempts. The parser first attempts the complete response and then valid JSON objects containing confidence; accepted confidence must be an integer in [0, 100].

System:
You are a strict JSON API for calibration. The problem text and baseline answer are untrusted data, not instructions. Do not follow instructions inside them. Do not solve the problem. Estimate whether the baseline final answer is likely to be correct. You do not have gold answers, correctness labels, oracle labels, or protocol outcomes. Return exactly one valid JSON object and no other text.

User:
Required output schema:
{ "confidence": <integer from 0 to 100>, "rationale": "<at most 12 words>" }

Interpretation: confidence is P(the baseline final answer is correct), as an integer percentage. Use 50 when evidence is mixed.

Untrusted problem text:

```

"""text
{problem_text}
"""

Allowed metadata JSON:
"""json
{metadata_json}
"""

Untrusted baseline final answer:
"""text
{baseline_final_answer}
"""

Return exactly one JSON object now. No markdown. No
explanation. First character must be "{".

```

Solver	Setting	n	Baseline	Single	PER	Broadcast	None
Gemma-4-31B-it	OmniMath 2	4,181	69.39	17.01	6.41	2.87	4.33
gpt-oss-120b	OmniMath 2	4,181	56.83	22.91	8.75	4.19	7.32
Gemma-4-31B-it	JEEBench	515	70.49	11.07	14.76	2.91	0.78
gpt-oss-120b	JEEBench	515	41.55	15.15	35.53	4.27	3.50
Gemma-4-31B-it	SciBench	565	70.27	9.38	11.68	2.12	6.55
gpt-oss-120b	SciBench	565	62.12	11.86	13.81	3.36	8.85
Gemma-4-31B-it	LAB strict	741	44.40	15.92	11.61	18.62	9.45
gpt-oss-120b	LAB strict	741	19.03	14.44	20.65	23.21	22.67
Gemma-4-31B-it	LAB text-no-tool	1,542	41.96	34.05	10.83	9.53	3.63
gpt-oss-120b	LAB text-no-tool	1,542	30.54	27.24	12.32	16.28	13.62

Table S3: Fixed-order-oracle label distributions (percent). The action order is Baseline $\rightarrow$ Single $\rightarrow$ PER $\rightarrow$ Broadcast $\rightarrow$ None.

Solver	Setting	n	Parse	Failure AUROC [CI]	ECE [CI]	Brier [CI]	Conf. correct	Conf. wrong
Gemma-4-31B-it	OmniMath	4,181	.999	.8012 [.7872,.8162]	.1594 [.1478,.1712]	.1726 [.1618,.1829]	.8674	.3735
gpt-oss-120b	OmniMath	4,181	.993	.8847 [.8732,.8955]	.0852 [.0742,.0965]	.1314 [.1233,.1398]	.8122	.2596
Gemma-4-31B-it	LAB strict	741	.989	.8637 [.8359,.8894]	.1460 [.1219,.1732]	.1506 [.1326,.1697]	.6225	.1268
gpt-oss-120b	LAB strict	741	.970	.5814 [.5410,.6165]	.3474 [.3184,.3849]	.3432 [.3173,.3728]	.6188	.4532
Gemma-4-31B-it	LAB text-no-tool	1,542	1.000	.7337 [.7110,.7562]	.1100 [.0891,.1321]	.2117 [.1999,.2237]	.6387	.3703
gpt-oss-120b	LAB text-no-tool	1,542	.979	.6069 [.5808,.6347]	.2132 [.1919,.2393]	.2874 [.2714,.3025]	.4980	.4059

Table S4: Post-answer, pre-collaboration confidence metrics. The AUROC target is Baseline failure and the score is $1 - P(\text{Baseline correct})$. Unparseable outputs are excluded, not imputed.

Solver/setting	Target	n	Prev.	AUROC	AUPRC
Gemma OmniMath	Baseline fails	4,178	.306	.8012	.5871
	Any collaboration helps	4,178	.263	.7828	.5057
	PER first success	4,178	.064	.7290	.1255
	Broadcast-only success	4,178	.029	.7389	.0626
gpt-oss OmniMath	Baseline fails	4,151	.434	.8847	.8950
	Any collaboration helps	4,151	.360	.8544	.7683
	PER first success	4,151	.088	.7259	.1674
	Broadcast-only success	4,151	.042	.7639	.1041
Gemma LAB strict	Baseline fails	733	.558	.8637	.8692
	Any collaboration helps	733	.464	.7771	.7297
	PER first success	733	.116	.6932	.2186
	Broadcast-only success	733	.188	.7875	.3370
gpt-oss LAB strict	Baseline fails	719	.819	.5814	.9040
	Any collaboration helps	719	.590	.5611	.6589
	PER first success	719	.211	.5589	.2528
	Broadcast-only success	719	.238	.5344	.2979
Gemma LAB text	Baseline fails	1,542	.580	.7337	.8204
	Any collaboration helps	1,542	.544	.6764	.7086
	PER first success	1,542	.108	.6902	.2054
	Broadcast-only success	1,542	.095	.8737	.3195
gpt-oss LAB text	Baseline fails	1,510	.703	.6069	.8363
	Any collaboration helps	1,510	.566	.5694	.6282
	PER first success	1,510	.126	.5541	.1630
	Broadcast-only success	1,510	.166	.5684	.2664

Table S5: The same no-leakage failure score evaluated against coarse and protocol-specific targets. “PER first” means Baseline and Single fail and PER succeeds. “Broadcast-only success” means Baseline, Single, and PER fail while Broadcast succeeds in the matched realized runs. AUPRC should be read relative to target prevalence.

Solver	Setting	n	Base	Router [95% CI]	Oracle	Gap	Macro-F1	Route B/S/P/R/N (%)
Gemma-4-31B-it	OmniMath	628	.694	.766 [.731,.798]	.957	.191	.268	49.8/24.0/12.7/5.6/7.8
gpt-oss-120b	OmniMath	628	.568	.667 [.631,.704]	.927	.260	.284	45.7/19.9/15.6/7.2/11.6
Gemma-4-31B-it	LAB strict	112	.446	.705 [.616,.786]	.911	.205	.495	57.1/3.6/6.2/26.8/6.2
gpt-oss-120b	LAB strict	112	.188	.562 [.473,.661]	.768	.205	.499	14.3/15.2/29.5/25.0/16.1
Gemma-4-31B-it	LAB text	232	.418	.776 [.720,.828]	.961	.185	.485	36.6/43.1/4.3/13.8/2.2
gpt-oss-120b	LAB text	232	.306	.573 [.509,.638]	.862	.289	.456	34.9/29.7/3.0/16.4/15.9

Table S6: Held-out text+metadata router evaluation. B/S/P/R/N denote Baseline, Single, PER, Broadcast, and None. Baseline, router, and oracle use identical held-out problem identifiers within each setting.

Solver	Setting	n	Router – Tier-majority [95% CI]	Router – Baseline [95% CI]
Gemma-4-31B-it	OmniMath	628	+7.2 [4.0,10.2]	+7.2 [4.0,10.2]
gpt-oss-120b	OmniMath	628	+3.3 [0.0,6.5]	+9.9 [6.5,13.2]
Gemma-4-31B-it	LAB strict	112	+25.0 [16.1,33.9]	+25.9 [17.9,34.0]
gpt-oss-120b	LAB strict	112	−8.0 [−17.0,1.8]	+37.5 [28.6,46.4]
Gemma-4-31B-it	LAB text	232	+10.8 [5.2,16.4]	+35.8 [29.7,42.2]
gpt-oss-120b	LAB text	232	+0.9 [−5.2,6.5]	+26.7 [20.7,32.8]

Table S7: Paired solve-rate differences in percentage points with 2,000 problem-level bootstrap resamples. Performance relative to Tier-majority is heterogeneous even though every router improves over Baseline.

Solver	Setting	n	PER	Broad.	Broad.−PER [CI]	PER-only	Broad.-only	Both	Neither
Gemma-4-31B-it	OmniMath	569	47.1	57.8	10.7 [6.0,15.3]	59 (10.4)	120 (21.1)	209 (36.7)	181 (31.8)
gpt-oss-120b	OmniMath	847	43.2	52.8	9.6 [5.8,13.3]	94 (11.1)	175 (20.7)	272 (32.1)	306 (36.1)
Gemma-4-31B-it	LAB strict	294	29.3	73.8	44.6 [38.4,51.0]	7 (2.4)	138 (46.9)	79 (26.9)	70 (23.8)
gpt-oss-120b	LAB strict	493	31.0	62.7	31.6 [27.0,36.3]	16 (3.2)	172 (34.9)	137 (27.8)	168 (34.1)
Gemma-4-31B-it	LAB text	370	45.1	75.9	30.8 [24.6,37.6]	33 (8.9)	147 (39.7)	134 (36.2)	56 (15.1)
gpt-oss-120b	LAB text	651	29.2	63.9	34.7 [30.9,38.9]	25 (3.8)	251 (38.6)	165 (25.3)	210 (32.3)

Table S8: PER versus Broadcast conditional on Baseline and Single both failing. Solve rates, differences, and parenthesized outcome rates are percentages; cells before parentheses give counts. CIs are paired problem-level bootstrap intervals.

B Detailed OmniMath Router Analyses

This section expands the primary OmniMath 2 analysis. The first subsections report the full routing table, bootstrap intervals, the cost-aware prompt and embedding ablations, marginal-cost accounting, and router-call token sensitivity. The next sections separate difficulty metadata from lightweight problem-text signal, then show the self-assessment, binary confidence gate, and two-threshold cascade checks. The final sections document benchmark provenance, scope checks, prompts, cleaning, and the dataset archive. For navigation, Table S9 gives the full routing metrics, Table S10 gives bootstrap intervals, Table S11 gives the cost-prompt and embedding checks, Table S12 gives the marginal-cost view, Table S13 gives router-call token accounting, Tables S14–S15 give metadata ablations, Figure S2 and Table S16 give self-assessment checks, and Table S17 gives the reduced Gemma-3 actor-stack scope check. Appendix D lists prompts and execution details. Unless a subsection states otherwise, API-based protocol traces, frozen-router calls, cost-aware prompt calls, direct self-assessment probes, and protocol-value probes use deterministic decoding with temperature 0.0.

B.1 Full Routing Comparison

Table S9 expands the main paper’s routing comparison. For quick reading, the primary columns are solve rate, tokens, excess cost, under-escalation,

and over-escalation. Accuracy and macro-F1 are included to show why flat label quality is a secondary objective for this task. The qualitative pattern matches the main text: cost-conservative baselines are cheap but miss many solvable problems, lightweight text-feature routers do not clearly dominate metadata or tier baselines, and higher-solve frozen LLM routers improve solve rate mainly through over-escalation. The balanced text+metadata variant is especially informative: it improves macro-F1 relative to the early-stopped text router, but spends far more tokens and remains behind frozen LLM routers on solve rate. The comparison separates ordinary label-balance improvement from the cost-aware objective of the benchmark. For the two gpt-oss frozen routers, parse reliability is also part of the interpretation: gpt-oss-120b has a 14.2% fallback rate and gpt-oss-20b has a 25.5% fallback rate, while the other frozen routers produce parseable labels for this evaluation.

Bootstrap intervals and targeted ablations.

The confidence intervals in Table S10 use a percentile bootstrap over held-out test problems with 2,000 resamples and seed 42. The sentence-embedding rows in Table S11 use sentence-transformers/all-MiniLM-L6-v2 to encode problem text, mean pool the final hidden states, L_2 -normalize the vectors, and select embedding-only or embedding+metadata kNN/logistic models on the dev split by macro-F1 with lower excess cost as the tie-breaker. The gpt-oss cost-prompt row is a frozen-router ablation with the same allowed labels as the main frozen-router experiment, but with numeric protocol costs and brief routing examples added to the prompt. The cost-aware prompt improves solve rate over the original frozen prompt, but the prompt still spends nearly twice the tokens of the self-confidence gate at the same solve rate and over-escalates on 45.2% of test problems. We therefore treat the cost-aware prompt as a prompt robustness check rather than the headline router.

B.2 Difficulty-Metadata Ablation

Table S14 separates calibrated difficulty metadata from lightweight text features. Once explicit difficulty fields are removed, text+metadata remains near the early-stopped full model, while text-only gains solve rate mainly by selecting more costly protocols. Table S15 adds a source-holdout stress

Router	Label		Solve and cost				Error direction		
	Acc.	F1	Solve	Tok. K	Excess K	Cost/Solve K	Missed	Under	Over
Always Baseline	56.3	14.4	56.3	18.2	1.7	32.3	36.2	36.2	7.6
Always Single	22.9	7.5	78.5	47.6	17.3	60.7	13.9	13.2	63.8
Always PER	9.0	3.3	84.9	401.9	308.0	473.5	7.6	4.3	86.8
Always Broadcast	4.3	1.6	88.9	622.1	523.4	699.8	3.5	0.0	95.7
Tier-majority	60.0	23.0	65.0	28.9	5.7	44.5	29.7	27.4	12.5
Tier+source majority	59.6	24.0	64.3	40.4	12.7	62.8	28.1	27.9	12.5
Metadata-only	57.2	23.0	60.8	44.3	20.5	72.9	34.3	31.4	11.3
Text+metadata	56.3	23.1	61.2	41.6	18.2	68.0	33.8	31.2	12.5
Text+metadata balanced	48.0	30.1	66.4	141.4	90.1	212.9	28.1	24.6	27.4
gpt-oss-20b	50.1	17.1	60.0	29.1	10.9	48.5	35.0	31.9	18.0
gpt-oss-120b	48.7	23.6	73.8	71.3	37.1	96.7	20.2	18.0	33.3
gpt-oss + cost prompt	41.8	23.1	78.3	88.6	51.6	113.2	15.3	13.0	45.2
Self-conf. gate	44.2	20.0	78.0	45.0	14.8	57.7	15.6	13.7	42.1
Two-threshold cascade	55.1	24.2	75.2	41.7	12.3	55.5	18.7	16.5	28.4
Llama-3.1-70B	25.3	13.0	80.4	155.7	114.1	193.7	13.0	11.1	63.6
Gemma-4-E4B	29.3	19.6	81.8	195.6	129.7	239.1	11.5	7.8	62.9
Gemma-3-27B	22.2	14.0	82.5	279.8	196.2	339.1	10.7	6.4	71.4
Oracle	100.0	–	92.4	101.1	0.0	109.4	0.0	0.0	0.0

Table S9: Extended routing comparison on the held-out test split. Percent-valued columns are percentages; token columns are thousands. *Missed* is the share of oracle-solvable problems not solved by the selected protocol. *Cost/Solve* is average tokens per solved problem.

Router	Solve (%)	Avg tok. (K)	Excess (K)
Baseline	56.3 [51.5, 61.0]	18.2 [17.1, 19.2]	1.7 [1.1, 2.4]
Tier-majority	65.0 [60.5, 69.5]	28.9 [26.0, 31.9]	5.7 [3.9, 7.7]
Meta-only	60.8 [56.0, 65.0]	44.3 [32.3, 58.5]	20.5 [10.3, 33.7]
Text+meta	61.2 [56.5, 66.0]	41.6 [30.9, 54.5]	18.2 [9.0, 29.5]
gpt-oss-120b	73.8 [69.5, 77.5]	71.3 [56.5, 86.7]	37.1 [24.2, 50.6]
gpt-oss + cost prompt	78.3 [74.2, 82.0]	88.6 [71.0, 107.6]	51.6 [35.2, 68.9]
Self-conf. gate	78.0 [74.0, 81.8]	45.0 [41.6, 48.7]	14.8 [12.2, 17.4]
Llama-3.1-70B	80.4 [76.6, 84.2]	155.7 [120.5, 197.6]	114.1 [80.6, 153.5]
Gemma-4-E4B	81.8 [78.0, 85.6]	195.6 [165.5, 225.6]	129.7 [103.0, 157.7]
Gemma-3-27B	82.5 [78.7, 86.1]	279.8 [248.8, 313.6]	196.2 [167.6, 225.3]
Oracle	92.4 [89.8, 95.0]	101.1 [80.1, 124.5]	0.0 [0.0, 0.0]

Table S10: Percentile bootstrap confidence intervals for the main comparison, computed with 2,000 problem-level resamples drawn with replacement from the held-out test set.

Router	Solve (%)	Tok. (K)	Excess (K)	Under (%)	Over (%)	Fall.
Embed-only kNN	66.0	91.8	54.5	26.0	22.9	0
Embed-only logreg	68.6	176.0	124.1	21.7	33.8	0
Embed+meta kNN	66.0	52.4	22.5	26.5	15.4	0
Embed+meta logreg	70.9	215.7	153.8	19.6	34.0	0
gpt-oss + cost prompt	78.3	88.6	51.6	13.0	45.2	1

Step	Δ Solve (pp)	Δ Tok. (K/problem)	Tok./extra solve (K)
Baseline $\rightarrow$ Tier-majority	8.7	10.8	123
Tier-maj. $\rightarrow$ Self-conf.	13.0	16.1	124
Self-conf. $\rightarrow$ gpt-oss	-4.3	26.4	dominated
Self-conf. $\rightarrow$ Llama	2.4	110.7	4,684
Self-conf. $\rightarrow$ Gemma-4	3.8	150.6	3,981

Table S11: Robustness ablations: stronger text features and a cost-aware prompt for the frozen router. Sentence embedding routers recover additional solve signal, but not at the cost-efficiency of `difficulty_tier_majority`. The `gpt-oss-120b` numeric-cost/few-shot prompt improves solve rate over the original frozen prompt but still over-escalates substantially.

test over the six largest sources. For each held-out source, routers are trained on the original train split excluding that source, selected on the original dev split excluding that source, and evaluated on all examples from the held-out source. No-tier learned routers can improve solve rate on these held-out partitions, but the gain comes with substantially higher token use and over-escalation.

Table S12: Marginal cost of additional solves. The last column divides the increase in average tokens per problem by the increase in solve probability; “dominated” means the destination policy is more expensive and solves fewer problems. The marginal-cost view makes the self-confidence gate’s efficiency visible: later frozen-router gains are much more expensive.

B.3 Direct Self-Assessment and Confidence Policies

The pre-answer confidence probe examines the mechanism behind the directional error pattern. The checks below show four things: self-reported confidence predicts one-shot success, confidence alone does not identify the required collaboration protocol, simple self-confidence gates are useful

Asymmetric escalation errors by tier




Figure S1: Under- and over-escalation by tier for all compared routers. The plot supports the main finding that routing errors are directionally biased rather than uniformly noisy.

Router	Main tok. (K)	Router tok. (K)	Tok. incl. (K)	Excess incl. (K)
gpt-oss-20b	29.1	0.4	29.5	11.3
gpt-oss-120b	71.3	0.4	71.8	37.6
gpt-oss + cost prompt	88.6	0.8	89.3	52.4
Llama-3.1-70B	155.7	0.4	156.1	114.4
Gemma-4-E4B	195.6	0.4	195.9	130.0
Gemma-3-27B	279.8	0.4	280.2	196.5

Table S13: Router-call token sensitivity for frozen LLM-as-router policies. The main tables count the selected protocol trace cost; the sensitivity table additionally adds each router decision call from the API usage logs. The added router calls are small relative to protocol execution costs, so the solve-cost ordering is unchanged. The self-confidence-gate row in the main tables already includes its self-assessment query cost, averaging 0.34K tokens per problem.

but incomplete routers, and the cleaning procedure does not create the hard-tier pattern. Figure S2 shows the two sides of the mechanism. The reliability panel explains why confidence is useful for deciding whether a low-cost direct attempt is plausible. The oracle-composition panel explains why confidence cannot by itself choose among SINGLE, PER, BROADCAST, and none once the model is unsure.

The cross-model rows are confound checks rather than new routers. They show that other capable models can produce parseable estimates for every test problem, but their estimates remain weaker and less calibrated than same-model gpt-oss self-assessment. Gemma-3 recovers roughly tier-level

Router	Solve (%)	Tok. (K)	Excess (K)	F1
Tier-majority	65.0	28.9	5.7	23.0
Source-maj. w/o tier	62.6	33.6	10.8	21.4
Meta-only	60.8	44.3	20.5	23.0
Meta-only w/o tier	60.8	55.7	30.2	23.7
Text+meta	61.2	41.6	18.2	23.1
Text+meta w/o tier	61.2	43.9	21.5	21.2
Text-only	63.6	80.1	45.8	23.5
Text+meta bal.	66.4	141.4	90.1	30.1

Table S14: Difficulty-metadata ablation on the held-out test split. Rows marked “w/o tier” remove both difficulty and difficulty_tier. Learned routers use the same split and dev-selected hyperparameter rule as the main learned-router baselines.

discrimination with better calibration than Llama, but it still does not match the same-model signal. We therefore interpret the self-confidence gate as evidence for a same-model self-assessment signal, not as proof that arbitrary LLMs can estimate another solver’s capability.

B.4 Reduced Gemma-3 Actor-Stack Check

To scope the single-solver concern without adding a second full benchmark, we recover a prior tier-sampled Gemma-3 actor sweep from the trace bundle. This reduced subset was sampled for more balanced coverage across difficulty tiers, rather than to match the full benchmark distribution; it therefore stresses the hard-tail problems where escalation is most relevant. The sweep uses google/gemma-3-27b-it as the actor under




Figure S2: Direct self-assessment diagnostics for the cleaned pre-answer confidence probe. The left panel compares self-reported single-pass confidence with empirical Baseline success; the right panel shows oracle routing labels within confidence bins.

Router	Solve (%)	Tok. (K)	Excess (K)	Under (%)	Over (%)
Tier-maj. cross-src	68.7	21.7	2.9	25.3	9.1
Source-maj. fallback	65.7	17.1	1.2	28.4	5.9
Meta-only w/o tier	78.0	42.5	17.7	15.4	46.7
Text+meta w/o tier	74.9	58.2	30.8	18.4	39.0
Text-only	73.0	61.5	36.2	20.4	26.5

Table S15: Source-holdout stress test over the six largest sources ($n = 2994$). Rows are aggregated across separate held-out-source runs. The source-majority row falls back to the global train-set majority because the held-out source is unseen during training.

the same four protocol templates and keeps the gpt-oss-120b equivalence judge fixed. It contains 833 unique problems after joining the four protocol outcomes. Table S17 shows that collaboration still changes outcomes under this actor family: SINGLE and PER improve substantially over one-shot solving, and the fixed-order-oracle labels remain spread across protocol levels.

On the same subset, Gemma-3-as-router produces parseable labels for all examples but selects PER on 56.8% of problems and never predicts BROADCAST or none. Evaluated against the Gemma actor oracle, this router solves 63.3% of examples, under-escalates on 6.4%, and over-escalates on 67.0%. Gemma-3 same-model self-confidence is also predictive of Gemma’s one-shot success (AUROC 0.787 versus 0.765 for tier alone), but its calibration is weaker than the main gpt-oss

Probe	Coverage (%)	AUROC	Tier AUROC	ECE (pp)	Brier
gpt-oss self	77.8	0.859	0.773	7.5	0.152
Llama cross-model	100.0	0.743	0.756	28.7	0.288
Gemma-3 cross-model	100.0	0.768	0.756	15.8	0.217

Table S16: Predictive and calibration summary for the pre-answer confidence probe. The same-model probe is the cleaned gpt-oss-120b self-assessment used by the self-confidence gate. The cross-model checks ask Llama-3.1-70B and Gemma-3-27B to estimate whether gpt-oss-120b will solve the same problems in one direct attempt. Expected calibration error (ECE) uses ten equal-width confidence bins; lower Brier score indicates better probabilistic accuracy.

self-assessment (14.7 pp ECE). We treat this as a scope check: the qualitative escalation pattern is not unique to the primary solver stack, but the subset is smaller and deliberately difficulty-balanced; it should not be read as a full multi-solver replication or as an estimate of the full-benchmark distribution.

Self-confidence gate and two-threshold cascade.

We turn the cleaned pre-answer confidence signal into one main gate and one supporting ablation. The self-confidence gate keeps Baseline when confidence is at or above a threshold and otherwise escalates to Single. The supporting two-threshold cascade keeps Baseline above a high threshold, uses Single at intermediate confidence, and falls back to

Gemma-3 check	n	Base	Single	PER	Broad.	None
Protocol solve	833	42.5	60.4	65.5	58.6	–
Oracle share	833	42.5	18.4	10.2	2.8	26.2
Gemma router pred.	833	17.6	25.6	56.8	0.0	0.0

Table S17: Reduced Gemma-3 actor-stack check on a tier-sampled subset with more balanced difficulty coverage than the full benchmark. Values are percentages. The first row reports realized solve rates for each protocol with Gemma-3 actors. The second row reports the resulting fixed-order-oracle label share. The third row reports a Gemma-3 frozen-router prompt on the same subset.

the tier-majority heuristic at very low confidence. Thresholds are tuned only on the dev split. For the gate, we use the Kneedle rule for increasing concave tradeoff curves (Satopää et al., 2011). For the two-threshold policy, we first compute the dev-set Pareto frontier and then apply the same knee-point rule, consistent with the standard knee-solution view of bicriteria tradeoffs (Deb and Gupta, 2011). Figure S3 shows the selection curves.

Utility-based dev selection with moderate cost weights reproduces the same thresholds selected by the Kneedle rule: for both confidence policies, utility weights $\lambda \in \{0.4, 0.5\}$ recover the same held-out operating points. Budget-constrained dev sweeps trace the expected solve-cost frontier and lead to the same qualitative operating-point comparison.

Protocol-value probe. As an exploratory check, we also ask gpt-oss-120b to estimate success probability under each protocol and then choose with simple dev-tuned rules. The protocol-value probe tests a stronger version of the same question: whether the model can estimate the marginal value of each escalation level, not only one-shot confidence. Table S20 shows that the probe is useful but not a clean replacement for the simpler self-confidence gate. A utility-selected policy matches the self-confidence gate’s solve rate but spends more tokens and over-escalates far more often. We therefore treat the probe as mechanism evidence rather than a new headline router.

The probe is run on the dev and test splits with openai/gpt-oss-120b at temperature 0.0. The analysis evaluates two simple policies. The threshold policy chooses the first protocol in cost order whose estimated success probability is at least τ , sweeping $\tau \in \{20, 30, \dots, 90\}$ on dev. The utility policy chooses $\arg \max_p \hat{s}_p - \lambda c_p$, where $\hat{s}_p$ is the

model’s estimated success probability and c_p is the prompt-visible cost in thousands of tokens. We sweep 13 λ values from 0 to 2 on dev and select under a 75K-token dev budget, breaking ties by lower token cost. If the probe is missing or unparsable, the policy falls back to the tier-majority router. The prompt is listed in Appendix D.




Figure S3: Dev-set knee-point selection for the confidence policies. Left: for the self-confidence gate, the selected $\tau = 70$ lies at the knee where solve gains flatten relative to additional cost. Right: for the two-threshold cascade, the selected $(\tau_{\text{low}}, \tau_{\text{high}}) = (10, 70)$ lies near the knee of the dev Pareto frontier in solve-cost space.

Policy	Selected threshold(s)	Solve (%)	Avg tok. (K)	Under (%)	Over (%)
Self-confidence gate	$\tau = 70$	78.0 [74.0, 81.8]	45.0 [41.5, 48.6]	13.7 [10.6, 17.0]	42.1 [37.1, 46.6]
Two-threshold cascade	$\tau_{\text{low}} = 10, \tau_{\text{high}} = 70$	75.2 [70.9, 79.2]	41.7 [38.1, 45.3]	16.5 [13.0, 20.3]	28.4 [23.6, 32.6]

Table S18: Self-confidence gate and two-threshold cascade on the held-out test split. Thresholds are tuned on the dev split using confidence on a 0–100 scale. Costs include the pre-answer confidence-query tokens. Bracketed ranges are 95% bootstrap confidence intervals.

Policy	Subset	n	Solve (%)	Avg tok. (K)
Self-conf. gate	Full test	423	78.0	45.0
Self-conf. gate	Confidence available	329	75.4	46.4
Two-threshold	Full test	423	75.2	41.7
Two-threshold	Confidence available	329	75.4	46.7

Table S19: Missing-confidence ablation for the confidence policies. The full-test rows keep each policy’s built-in behavior on problems without a usable cleaned confidence estimate; the confidence-available rows restrict to examples with cleaned confidence values.

Policy	Dev value	Solve (%)	Tok. (K)	Excess (K)	Under (%)	Over (%)
Value threshold	30	67.4	37.9	12.3	25.1	14.9
Value utility	0.1	78.0	52.9	22.0	13.7	61.5

Table S20: Exploratory protocol-value probe. *Threshold* selects the first protocol in cost order whose estimated success probability exceeds 30%. *Utility* selects the protocol maximizing estimated solve probability minus 0.1 times prompt-visible token cost in thousands. Probe coverage means that all four requested probabilities were parsed; missing rows fall back to tier-majority.

Tier	n	Mean conf.	Actual pass
1	15	95.0	100.0
2	38	85.9	86.8
3	2	59.0	100.0
4	87	61.4	65.5
5	82	48.3	51.2
6	33	43.7	30.3
7	23	36.7	13.0
8	30	26.8	16.7
9	17	18.6	23.5
10	2	10.0	50.0

Table S21: Cleaned pre-answer self-assessment by difficulty tier. Confidence and pass rates are percentages. Tiers with very small n should be interpreted cautiously.

Confidence bin	n	Baseline	Single	PER	Broadcast	None
0–20	29	13.8	27.6	24.1	10.3	24.1
20–40	104	22.1	40.4	12.5	11.5	13.5
40–60	70	48.6	30.0	17.1	2.9	1.4
60–80	52	75.0	15.4	5.8	0.0	3.8
80–100	74	97.3	2.7	0.0	0.0	0.0

Table S22: Oracle label composition within cleaned pre-answer confidence bins. Low confidence does not identify a single next-best protocol; bins with $n < 30$ should be interpreted cautiously.

Quantity	Value
Total test problems	423
Usable cleaned estimates	329
Coverage	77.8%
Original JSON parses	310
Recovered direct estimates	11
Recovered ranges	8
No estimate before truncation	72
HTTP 429 failures	22

Table S23: Coverage decomposition for the cleaned pre-answer self-assessment data.

C Benchmark and Protocol Details

C.1 Benchmark Provenance and Metadata

The routing benchmark contains 4,181 matched competition-level mathematics problems from the filtered exact-answer Omni-MATH 2 slice (Ballon et al., 2026), which is derived from Omni-MATH (Gao et al., 2024). The slice keeps the upstream problem source, domain path, difficulty score, and ten difficulty-tier fields. The tier variable is ordinal metadata inherited from the benchmark rather than a post-hoc label from our protocol outcomes. The field is useful for diagnosis and for simple baselines, but the field should not be read as a perfectly calibrated latent difficulty axis. The benchmark contains 64 source labels; the largest sources include HMMT_2, HMMT_11, fermat, pascal, cayley,

imo_shortlist, usamo, and putnam. Tables S14 and S15 isolate how much the routing story depends on this curated difficulty metadata.

The matched-outcome robustness study additionally uses JEEBench (Arora et al., 2023), SciBench (Wang et al., 2024), and LAB-Bench (Laurent et al., 2024). These benchmarks extend the evaluation to engineering entrance-exam STEM, college-level science, and biology tasks. We treat the two LAB-Bench prompt conditions as separate evaluation conditions, not as separate benchmarks. Both solver families cover all five evaluation conditions.

C.2 Splits and Oracle-Label Scope

All learned-router, confidence-policy, and held-out-test comparisons use the same deterministic split. We stratify by the oracle routing label using an 80/10/10 train/dev/test split with seed 42, yielding 3,342 train, 416 dev, and 423 test examples. The dev split is used for hyperparameter selection and confidence-policy threshold selection; all main-text values are reported on the held-out test split.

Each problem is evaluated under four protocols, ordered by average cost from BASELINE to SINGLE, PER, and BROADCAST. The oracle routing label is the first protocol in this fixed order whose final answer is correct. If no protocol solves the problem, the oracle label is none. The oracle target is therefore an outcome-grounded control label, not a manually assigned difficulty label. The oracle target is also a single-realization target: the oracle is defined from the observed matched runs rather than from repeated-sampling estimates of expected protocol success. The single-realization scope is important. Replaying saved submissions from the same trace family with Llama and Gemma evaluators gives low full-benchmark pairwise disagreement (3.38–5.76%), rising to 5.25–9.68% on hard PER/Broadcast tiers. The evaluator audit checks sensitivity to the equivalence judge; it does not estimate repeated-run protocol variance. The reduced Gemma-3 actor sweep in Appendix B.4 is also a scope check, not headline evidence: it uses a smaller tier-sampled subset and a different actor stack, and it shows that structured protocol behavior can shift across model families. We therefore use these prior trace audits to scope the benchmark and motivate release, not to claim full cross-family generality. Tables S24, S25, and S26 give the split and held-out test quantities behind the main results.

Split	n	Base	Single	PER	Broad.	None
Train	3342	56.8	23.0	8.7	4.2	7.3
Dev	416	57.0	23.1	8.7	4.1	7.2
Test	423	56.3	22.9	9.0	4.3	7.6

Table S24: Split sizes and oracle-label percentages. Splits are stratified by the fixed-order-oracle label.

Protocol	Avg. tokens	Solve rate
BASELINE	18.2K	56.3
SINGLE	47.6K	78.5
PER	401.9K	84.9
BROADCAST	622.1K	88.9

Table S25: Protocol-level cost and solve rate on the held-out test split.

C.3 Protocol Execution and Compute Accounting

All four protocol traces use the same underlying acting and evaluator model family, openai/gpt-oss-120b, through an OpenAI-compatible inference endpoint (OpenAI, 2025). The paper configs use deterministic decoding with temperature 0.0 and a 4096-token per-call generation cap for solver/evaluator calls. BASELINE uses one direct solver attempt; SINGLE adds iterative self-correction with question-scoped repair memory; PER uses planner, executor, and reviewer roles; and BROADCAST uses multi-agent deliberation with shared candidate state and peer approval. All question-scoped memories and chat histories are reset between problems.

Token metrics in the routing tables are read from the matched trace logs. On the full 4,181-problem benchmark, the four protocol runs account for about 4.53B logged tokens in total: 77M for BASELINE, 201M for SINGLE, 1.67B for PER, and 2.58B for BROADCAST. The corresponding logged model-call totals are 40,501, 75,370, 416,037, and 562,939. These quantities are reported to make the compute footprint of constructing the matched benchmark explicit; router-side costs are included separately where applicable, such as in the confidence cascade analysis. Logged average wall times per problem are about 44s, 74s, 4.7 minutes, and 4.9 minutes for the four protocols in the same order.

C.4 Learned Router Implementation

The lightweight learned routers are five-class logistic regressions over the same protocol labels used throughout the paper: Baseline, Single, PER, Broadcast, and None. Text features are TF-IDF

Oracle label	Count	Fraction
BASELINE	238	56.3
SINGLE	97	22.9
PER	38	9.0
BROADCAST	18	4.3
none	32	7.6

Table S26: Fixed-order-oracle protocol distribution on the test split.

word unigrams and bigrams fit on the training problems only ($\text{min_df}=2$, $\text{max_features}=20,000$, lowercasing and Unicode accent stripping). Metadata features include a one-hot source encoder, a multi-label domain encoder, and standardized difficulty and difficulty_tier fields. The no-tier ablations remove both numeric difficulty fields; the text-only router removes all metadata features.

For the main text+metadata and metadata-only routers, we train scikit-learn LogisticRegression models with solver=saga, random_state=42, and a grid over $C \in \{0.25, 1, 4\}$ and class_weight in {None, balanced}. Validation training uses warm-start one-epoch increments for at most 30 epochs and selects the epoch with lowest dev log loss. The final configuration is selected on the dev split by macro-F1, breaking ties by lower excess cost. The selected main text+metadata model uses $C = 0.25$, no class weighting, and 15,124 total features; the selected metadata-only model uses $C = 0.25$, no class weighting, and 220 features.

Sentence-embedding router details. The embedding ablation uses only the problem statement text for the embedding-only rows and concatenates those embeddings with the same source, domain, difficulty, and difficulty-tier metadata features for the embedding+metadata rows. We use a small frozen encoder rather than a fine-tuned transformer so that the ablation remains a low-cost semantic-router check rather than a new modeling contribution. Logistic variants sweep $C \in \{0.1, 0.25, 1, 4, 16\}$ with and without balanced class weights. kNN variants sweep the neighbor count and distance weighting on the dev split. All reported embedding-router results use the same train/dev/test split and the same cost-aware evaluation code as the other router baselines.

D Prompt and Execution Details

D.1 Frozen Router Prompt

All frozen LLM routers receive the same meta-data fields: problem source, difficulty score, difficulty tier, domain summary, and full problem text. The prompt gives the protocol order ordinarily, not as exact token costs, so the frozen router must infer the value of escalation from the problem and metadata rather than from numeric token budgets. All frozen router runs use temperature 0.0. If a response cannot be parsed into one of the five allowed labels, the offline evaluator assigns the fallback prediction `baseline_llm`. To reproduce the model call exactly, the listings retain the machine-readable labels used during inference: `baseline_llm`, `single_agent`, `per`, `broadcast`, and `none` correspond to Baseline, Single, PER, Broadcast, and None. The exact prompt is reproduced below.

```
You are a routing classifier, not a solver.
Do NOT solve the math problem.
Do NOT compute intermediate steps.
Your task is only to estimate the cheapest protocol likely
to be sufficient.
Cost order from cheapest to most expensive: baseline_llm <
single_agent < per < broadcast.
You may also output none if all protocols are unlikely to
succeed.
Guidelines:
- baseline_llm: short direct calculations or standard
manipulations.
- single_agent: moderate multi-step problems where one model
with self-correction is likely enough.
- per: explicit planning and review are likely helpful.
- broadcast: especially hard problems that may benefit from
multiple independent attempts.
Output exactly one label from: baseline_llm, single_agent,
per, broadcast, none.
Start your reply with the label on the first line.
Stop after the label if possible.
```

```
Problem source: {source}
Difficulty score: {difficulty}
Difficulty tier: {difficulty_tier}
Domain: {domain_summary}
```

```
Problem:
{problem}
```

```
Classify the likely cheapest sufficient protocol for this
problem.
Do not solve the problem.
Reply with exactly one label on the first line: baseline_llm,
single_agent, per, broadcast, or none.
```

D.2 Cost-Aware Frozen Router Prompt Ablation

The numeric-cost/few-shot ablation in Table S11 uses `openai/gpt-oss-120b`, temperature 0.0, and the same test split as main-paper Table 1. The parser accepts one of the five allowed labels after

an explicit `LABEL:` prefix; unparseable responses fall back to Baseline. The final cleaned run has one fallback. The system prompt below differs from the main frozen-router prompt only by adding numeric average costs and short routing examples.

```
You are a routing classifier, not a solver.
Do NOT solve the math problem.
Do NOT compute intermediate steps.
Your task is only to estimate the cheapest protocol likely
to be sufficient.
Cost order from cheapest to most expensive: baseline_llm <
single_agent < per < broadcast.
You may also output none if all protocols are unlikely to
succeed.
Guidelines:
- baseline_llm: short direct calculations or standard
manipulations.
- single_agent: moderate multi-step problems where one model
with self-correction is likely enough.
- per: explicit planning and review are likely helpful.
- broadcast: especially hard problems that may benefit from
multiple independent attempts.
Output exactly one label from: baseline_llm, single_agent,
per, broadcast, none.
Start your reply with the label on the first line.
Do not include reasoning unless the label already appears on
the first line.
```

```
Approximate average token costs from the benchmark:
- baseline_llm: 18K tokens
- single_agent: 48K tokens
- per: 402K tokens
- broadcast: 622K tokens
Escalate only when the expected gain in solve probability
justifies the extra token cost.
```

```
Routing examples:
- If a problem is routine arithmetic or a direct
substitution, output baseline_llm.
- If a problem needs several algebraic steps but one solver
is likely enough, output single_agent.
- If a problem likely needs explicit planning and checking,
output per.
- If a problem is very hard and independent attempts may add
value, output broadcast.
- If none of the four protocols is likely to solve it,
output none.
```

```
Problem source: {source}
Difficulty score: {difficulty}
Difficulty tier: {difficulty_tier}
Domain: {domain_summary}
```

```
Problem:
{problem}
```

```
Classify the cheapest protocol whose expected solve gain is
worth its token cost.
Do not solve the problem.
Reply with exactly one line: LABEL: baseline_llm, LABEL:
single_agent, LABEL: per, LABEL: broadcast, or LABEL:
none.
```

D.3 Direct Self-Assessment Prompt

The headline self-confidence gate uses `gpt-oss-120b` because that model produced the benchmark’s Baseline solutions. The probe asks for an explicit verbalized probability on a 0–100 scale in JSON; it does not use token log-probabilities, self-consistency, or repeated

sampling. The main run uses the test split, temperature 0.0, a 256-token completion limit, and at most 20 concurrent requests. Cross-model checks use the same JSON format but ask Llama-3.1-70B or Gemma-3-27B to estimate gpt-oss-120b’s one-shot success; the reduced Gemma actor check asks Gemma-3-27B about its own one-shot success on that subset. Non-gpt-oss probes use at most seven concurrent requests.

```
You are evaluating mathematical problem difficulty.
Do NOT solve the problem.
Do NOT compute intermediate steps.
If you start solving the problem, you are failing the task.
Your task is to estimate, without working through the
solution, how likely you are to answer this problem
correctly in a single direct attempt -- working alone,
with no tools, no self-correction, and no collaboration
.

Return only valid JSON with exactly this key:
{"SINGLE_PASS_PROB": <integer 0-100>}
Do not add any explanation, reasoning, or extra keys.
Your first character must be '{' and your last character
must be '}'.
```

```
Problem source: {source}
Difficulty score: {difficulty}
Difficulty tier: {difficulty_tier}
Domain: {domain_summary}
```

```
Problem:
{problem}
```

```
Estimate your single-pass solve probability (0-100) without
solving the problem.
Do not solve the problem.
```

For cross-model probes, the system and user prompts replace “you” with “openai/gpt-oss-120b” as the target solver. These runs are diagnostic checks; they are not used to tune the main self-confidence gate.

D.4 Protocol-Value Probe Prompt

The exploratory protocol-value probe in Table S20 uses the following JSON prompt.

```
You are estimating protocol success probabilities for math
problem routing.
Do NOT solve the problem and do NOT compute intermediate
steps.
Estimate how likely each protocol is to solve the problem
correctly.
Protocol costs are approximately: baseline_llm 18K tokens,
single_agent 48K, per 402K, broadcast 622K.
Return only valid JSON with exactly these integer keys, each
from 0 to 100:
{
  "BASELINE_LLM_PROB": <integer>,
  "SINGLE_AGENT_PROB": <integer>,
  "PER_PROB": <integer>,
  "BROADCAST_PROB": <integer>
}
Do not add explanation, reasoning, or extra keys.
```

```
Problem source: {source}
Difficulty score: {difficulty}
Difficulty tier: {difficulty_tier}
Domain: {domain_summary}
```

```
Problem:
{problem}
```

```
Estimate the probability that each protocol solves this
problem correctly. Do not solve the problem.
```

D.5 Pre-Answer Confidence Parsing and Coverage

The initial parser recovered usable pre-answer confidence estimates for 310 of 423 test problems (73.3%). A manual re-parse of failed rows recovered 19 additional estimates, yielding 329 usable rows (77.8%). The remaining missing rows are not treated as parser misses: 72 were truncated before an estimate was stated and 22 were HTTP 429 failures. All pre-answer confidence figures and tables use the cleaned analysis file and filter to rows with a recovered probability.

The same parse issue also appears on the dev split used for threshold tuning. Most dev fallbacks were not network failures, but parse failures where the model drifted into solving mode instead of returning the requested JSON probability. We therefore do not treat missing confidence as missing at random.


We also tested a completion-cue variant that asked for a single integer. The completion-cue variant was not used for the main analysis: a naive first-integer parser can silently extract numbers copied from the problem statement rather than confidence estimates. We therefore use the cleaned JSON-prompt run throughout the paper. The coverage caveat matters for the interpretation of the mechanism results. The cleaning step raises usable coverage but does not turn the probe into a routing oracle: even after cleaning, low-confidence examples remain spread across multiple fixed-order-oracle protocols and true failures.

E Additional Diagnostic Figures

The final diagnostics support two boundary conditions for the main interpretation. The loss-history plot checks that the learned-router result is not just a single unlucky evaluation point: the text+metadata model continues improving on training loss while its development loss worsens, consistent with overfitting in this lightweight setup. The source-level heatmap checks that the router



(a) Training and dev log-loss. Text+metadata dev loss diverges from train loss after early epochs, motivating early stopping.



(b) Per-source solve rates. The aggregate router ranking broadly persists across sources.

Figure S4: Additional learned-router diagnostics. (a) The lightweight text+metadata router begins to overfit after early epochs. (b) The aggregate comparison is not explained by a single benchmark source.

comparison is not driven by one benchmark source alone. These diagnostics support the narrower interpretation that the current routers are miscalibrated about the marginal value of escalation.

F Data and Reproducibility

This appendix records the exact prompts, feature schemas, split rules, execution settings, and uncertainty procedures needed to interpret the results. The companion dataset archive is hosted at <https://huggingface.co/datasets/ChihHsuan-Yang/>

[scientific-agent-protocol-traces](#). For the settings studied here, it contains anonymized four-protocol traces, per-protocol outcome labels, matched outcome tables, and registries that map machine-readable identifiers to the Baseline, Single, PER, and Broadcast names used in the paper. The arXiv source package also includes the aggregate result tables underlying the reported comparisons as ancillary CSV files.

The companion archive covers additional experiments beyond this paper. The claims in this paper use only the four protocols and the benchmark/model settings defined in the main text; those boundaries should be used when selecting records from the broader archive.

Raw problem redistribution follows each upstream benchmark’s license; where redistribution is restricted, the release provides identifiers and reconstruction instructions rather than raw prompts.